

Activity_Course 3 Waze project lab

September 29, 2023

1 Waze Project

Course 3 - Go Beyond the Numbers: Translate Data into Insights Creator: Thang Huynh

Your team is still in the early stages of their user churn project. So far, you've completed a project proposal and used Python to inspect and organize Waze's user data.

You check your inbox and notice a new message from Chidi Ga, your team's Senior Data Analyst. Chidi is pleased with the work you have already completed and requests your assistance with exploratory data analysis (EDA) and further data visualization. Harriet Hadzic, Waze's Director of Data Analysis, will want to review a Python notebook that shows your data exploration and visualization.

A notebook was structured and prepared to help you in this project. Please complete the following questions and prepare an executive summary.

2 Course 3 End-of-course project: Exploratory data analysis

In this activity, you will examine data provided and prepare it for analysis.

The purpose of this project is to conduct exploratory data analysis (EDA) on a provided dataset.

The goal is to continue the examination of the data that you began in the previous Course, adding relevant visualizations that help communicate the story that the data tells.

This activity has 4 parts:

Part 1: Imports, links, and loading

Part 2: Data Exploration * Data cleaning

Part 3: Building visualizations

Part 4: Evaluating and sharing results

Follow the instructions and answer the question below to complete the activity. Then, you will complete an executive summary using the questions listed on the [PACE Strategy Document](#).

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

3 Visualize a story in Python

4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

4.1.1 Task 1. Imports and data loading

For EDA of the data, import the data and packages that will be most helpful, such as pandas, numpy, and matplotlib.

```
[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Read in the data and store it as a dataframe object called df.

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[4]: # Load the dataset into a dataframe
df = pd.read_csv('waze_dataset.csv')
```

```
[5]: df.isnull().sum()
```

```
[5]: ID                0
label                700
sessions             0
drives               0
total_sessions       0
n_days_after_onboarding  0
total_navigations_fav1  0
total_navigations_fav2  0
driven_km_drives      0
duration_minutes_drives  0
activity_days         0
driving_days          0
device               0
```

`dtype: int64`

4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document and those below where applicable to complete your code: 1. Does the data need to be restructured or converted into usable formats?

2. Are there any variables that have missing data?

1. Data does not need to be restructured or converted because it is already in a structured format
2. From previous code, we can see that 'label' has 700 missing rows

4.2.1 Task 2. Data exploration and cleaning

Consider the following questions:

1. Given the scenario, which data columns are most applicable?
 2. Which data columns can you eliminate, knowing they won't solve your problem scenario?
 3. How would you check for missing data? And how would you handle missing data (if any)?
 4. How would you check for outliers? And how would handle outliers (if any)?
1. 'label' appears to be the most applicable column since we are interested in finding out user churn rate. Other variables will also be helpful since they are all related to the churn rate.
 2. 'ID' column can be eliminated since they are just numbers in order and has nothing to do with determining user churn rate.
 3. We can use `df.info()` to see the non null columns, then the difference between all data rows and non null rows of each column will be the number of missing data for each variable. If missing data is due to randomness, we can still go ahead and perform further analysis. If that is not the case, we need to investigate the real cause of that missingness to have appropriate actions.
 4. If we know that outliers are either mistakes, typos or errors, we can delete outliers though it is not good practise to delete outliers. If dataset is small, we can reassign new values to replace outliers values. If models based on dataset are resitant to outliers or we just need to to EDA on the dataset, we are more likely to leave them in.

Data overview and summary statistics Use the following methods and attributes on the dataframe:

- `head()`
- `size`
- `describe()`
- `info()`

It's always helpful to have this information at the beginning of a project, where you can always refer back to if needed.

```
[6]: df.head(10)
```

```
[6]:   ID    label  sessions  drives  total_sessions  n_days_after_onboarding  \
0    0  retained      283     226      296.748273                2276
1    1  retained      133     107      326.896596                1225
2    2  retained      114      95      135.522926                2651
3    3  retained       49      40       67.589221                 15
4    4  retained       84      68      168.247020                1562
5    5  retained      113     103      279.544437                2637
6    6  retained        3        2      236.725314                 360
7    7  retained       39      35      176.072845                2999
8    8  retained       57      46      183.532018                 424
9    9   churned       84      68      244.802115                2997

      total_navigations_fav1  total_navigations_fav2  driven_km_drives  \
0                        208                        0      2628.845068
1                         19                        64     13715.920550
2                          0                        0      3059.148818
3                        322                         7       913.591123
4                        166                         5      3950.202008
5                          0                        0       901.238699
6                        185                        18      5249.172828
7                          0                        0      7892.052468
8                          0                        26      2651.709764
9                         72                         0      6043.460295

      duration_minutes_drives  activity_days  driving_days  device
0          1985.775061           28           19  Android
1          3160.472914           13           11   iPhone
2          1610.735904           14            8  Android
3           587.196542            7            3   iPhone
4          1219.555924           27           18  Android
5           439.101397           15           11   iPhone
6           726.577205           28           23   iPhone
7          2466.981741           22           20   iPhone
8          1594.342984           25           20  Android
9          2341.838528            7            3   iPhone
```

```
[3]: df.size
```

```
[3]: 194987
```

Generate summary statistics using the `describe()` method.

```
[5]: df.describe()
```

```
[5]:
```

	ID	sessions	drives	total_sessions	\
count	14999.000000	14999.000000	14999.000000	14999.000000	
mean	7499.000000	80.633776	67.281152	189.964447	
std	4329.982679	80.699065	65.913872	136.405128	
min	0.000000	0.000000	0.000000	0.220211	
25%	3749.500000	23.000000	20.000000	90.661156	
50%	7499.000000	56.000000	48.000000	159.568115	
75%	11248.500000	112.000000	93.000000	254.192341	
max	14998.000000	743.000000	596.000000	1216.154633	

	n_days_after_onboarding	total_navigations_fav1	\
count	14999.000000	14999.000000	
mean	1749.837789	121.605974	
std	1008.513876	148.121544	
min	4.000000	0.000000	
25%	878.000000	9.000000	
50%	1741.000000	71.000000	
75%	2623.500000	178.000000	
max	3500.000000	1236.000000	

	total_navigations_fav2	driven_km_drives	duration_minutes_drives	\
count	14999.000000	14999.000000	14999.000000	
mean	29.672512	4039.340921	1860.976012	
std	45.394651	2502.149334	1446.702288	
min	0.000000	60.441250	18.282082	
25%	0.000000	2212.600607	835.996260	
50%	9.000000	3493.858085	1478.249859	
75%	43.000000	5289.861262	2464.362632	
max	415.000000	21183.401890	15851.727160	

	activity_days	driving_days
count	14999.000000	14999.000000
mean	15.537102	12.179879
std	9.004655	7.824036
min	0.000000	0.000000
25%	8.000000	5.000000
50%	16.000000	12.000000
75%	23.000000	19.000000
max	31.000000	30.000000

And summary information using the `info()` method.

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
#   ...
```

```

---  -----
0    ID                14999 non-null  int64
1    label             14299 non-null  object
2    sessions          14999 non-null  int64
3    drives            14999 non-null  int64
4    total_sessions    14999 non-null  float64
5    n_days_after_onboarding 14999 non-null  int64
6    total_navigations_fav1 14999 non-null  int64
7    total_navigations_fav2 14999 non-null  int64
8    driven_km_drives    14999 non-null  float64
9    duration_minutes_drives 14999 non-null  float64
10   activity_days      14999 non-null  int64
11   driving_days       14999 non-null  int64
12   device             14999 non-null  object
dtypes: float64(3), int64(8), object(2)
memory usage: 1.5+ MB

```

4.3 PACE: Construct

Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

Consider the following questions as you prepare to deal with outliers:

1. What are some ways to identify outliers?
 2. How do you make the decision to keep or exclude outliers from any future models?
1. We can use boxplot to show the distribution of data then can detect outliers or we can also use mean() and median() to understand the range of data and identify outliers.
 2. If we know that outliers are either mistakes, typos or errors, we can delete outliers though it is not good practise to delete outliers. If dataset is small, we can reassign new values to replace outliers values. If models based on dataset are resitant to outliers or we just need to to EDA on the dataset, we are more likely to leave them in.

4.3.1 Task 3a. Visualizations

Select data visualization types that will help you understand and explain the data.

Now that you know which data columns you'll use, it is time to decide which data visualization makes the most sense for EDA of the Waze dataset.

Question: What type of data visualization(s) will be most helpful?

- Line graph
- Bar chart
- Box plot
- Histogram
- Heat map
- Scatter plot
- A geographic map


==> ENTER YOUR RESPONSE HERE

Begin by examining the spread and distribution of important variables using box plots and histograms.


sessions *The number of occurrence of a user opening the app during the month*

```
[30]: # Box plot
sns.boxplot(x=df['sessions'])
plt.title('sessions box plot')
```

```
[30]: Text(0.5, 1.0, 'sessions box plot')
```



```
[63]: # Histogram
sns.histplot(x=df['sessions'])
median = df['sessions'].median()
plt.axvline(median,color='red',linestyle='-')
plt.text(100,1000,'median=48.0', color='red')
plt.title('sessions box plot');
```

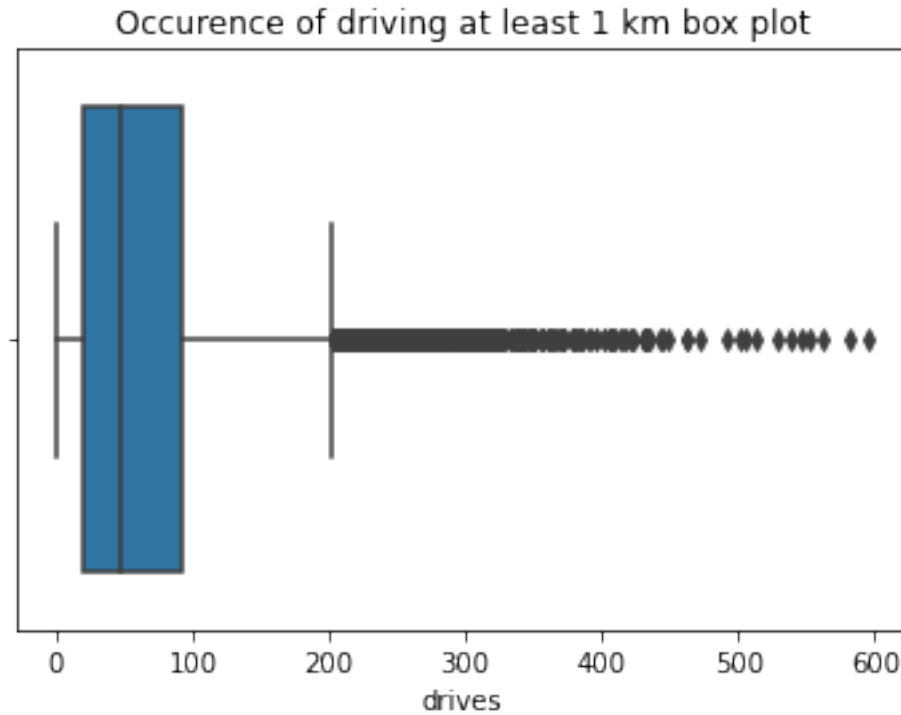


The `sessions` variable is a right-skewed distribution with half of the observations having 56 or fewer sessions. However, as indicated by the boxplot, some users have more than 700.

drives *An occurrence of driving at least 1 km during the month*

```
[68]: # Box plot
sns.boxplot(x=df['drives'])
plt.title('Occurence of driving at least 1 km box plot')
```

```
[68]: Text(0.5, 1.0, 'Occurence of driving at least 1 km box plot')
```

5 Histogram

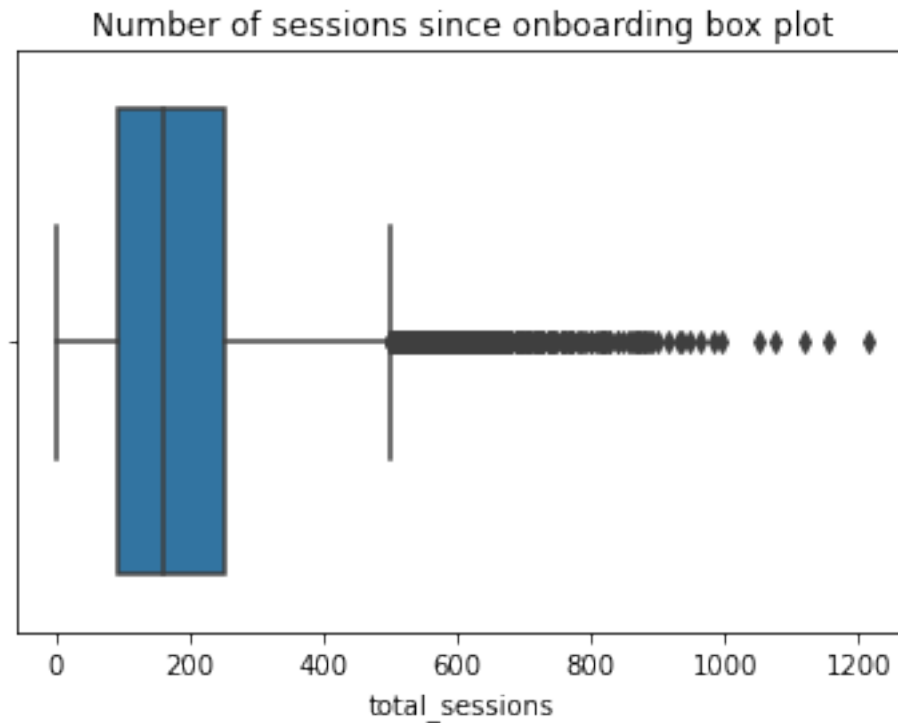
```
sns.histplot(df['drives']) median = df['drives'].median() plt.axvline(median, color='red', linestyle='-')
plt.text(60,1000,'median=', color='red') plt.title('Occurence of driving at least 1 km Histogram')
```

The **drives** information follows a distribution similar to the **sessions** variable. It is right-skewed, approximately log-normal, with a median of 48. However, some drivers had over 400 drives in the last month.

total_sessions *A model estimate of the total number of sessions since a user has onboarded*

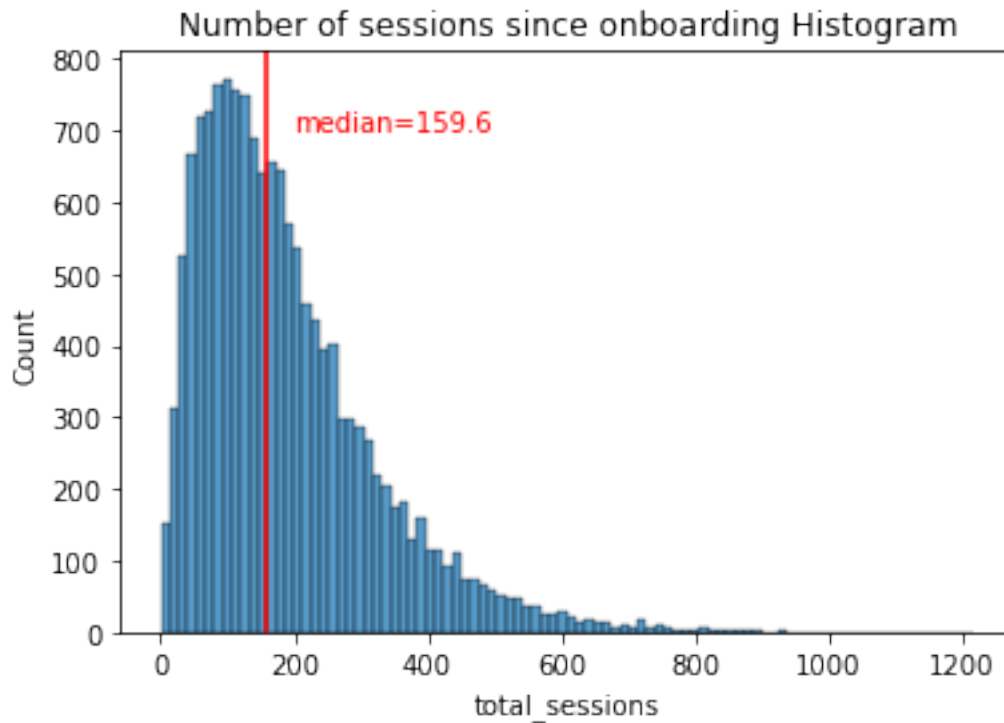
```
[67]: # Box plot
sns.boxplot(x=df['total_sessions'])
plt.title('Number of sessions since onboarding box plot')
```

```
[67]: Text(0.5, 1.0, 'Number of sessions since onboarding box plot')
```



```
[64]: # Histogram
sns.histplot(x=df['total_sessions'])
median = df['total_sessions'].median()
plt.axvline(median, color='red', linestyle='-')
plt.text(200,700,'median=159.6', color='red')
plt.title('Number of sessions since onboarding Histogram')
```

```
[64]: Text(0.5, 1.0, 'Number of sessions since onboarding Histogram')
```

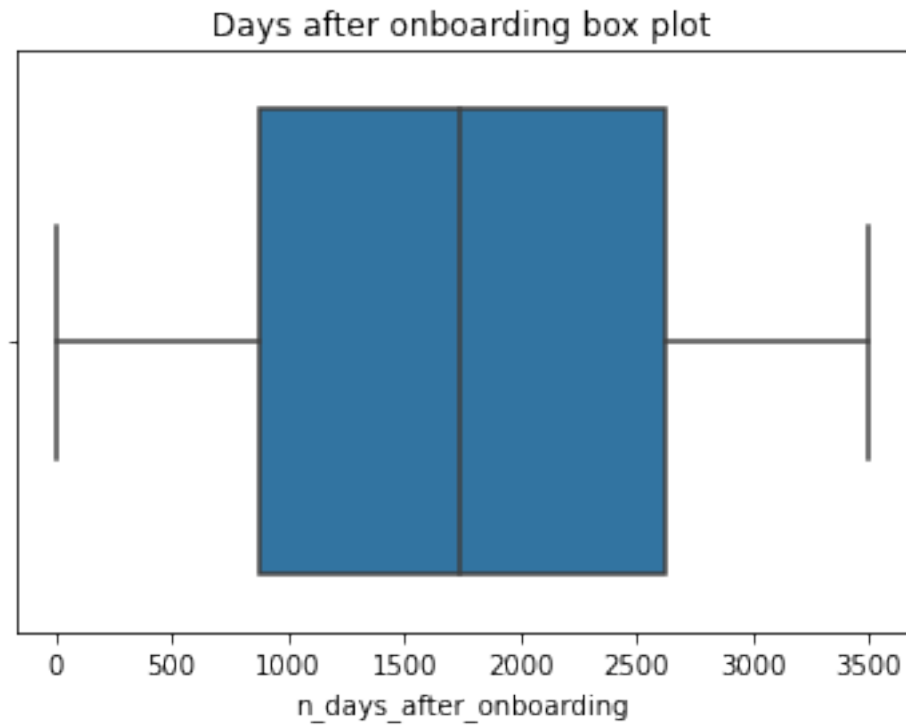


The `total_sessions` is a right-skewed distribution. The median total number of sessions is 159.6. This is interesting information because, if the median number of sessions in the last month was 48 and the median total sessions was ~160, then it seems that a large proportion of a user's total drives might have taken place in the last month. This is something you can examine more closely later.

`n_days_after_onboarding` *The number of days since a user signed up for the app*


```
[66]: # Box plot
sns.boxplot(x=df['n_days_after_onboarding'])
plt.title('Days after onboarding box plot')
```

```
[66]: Text(0.5, 1.0, 'Days after onboarding box plot')
```



```
[74]: # Histogram
sns.histplot(x=df['n_days_after_onboarding'])
median = df['n_days_after_onboarding'].median()
plt.axvline(median, color='red', linestyle='-')
plt.text(1750,650,'median=1741.0', color='red')
plt.title('Number of days since onboarding Histogram')
```

```
[74]: Text(0.5, 1.0, 'Number of days since onboarding Histogram')
```




The total user tenure (i.e., number of days since onboarding) is a uniform distribution with values ranging from near-zero to ~3,500 (~9.5 years).

driven_km_drives *Total kilometers driven during the month*


```
[75]: # Box plot
sns.boxplot(x=df['driven_km_drives'])
plt.title('driven km box plot')
```

```
[75]: Text(0.5, 1.0, 'driven km box plot')
```



```
[7]: # Histogram
sns.histplot(x=df['driven_km_drives'])
median = df['driven_km_drives'].median()
plt.axvline(median, color='red', linestyle='-')
plt.text(4000,700,'median=3493.9.0', color='red')
plt.title('driven km Histogram')
```

```
[7]: Text(0.5, 1.0, 'driven km Histogram')
```

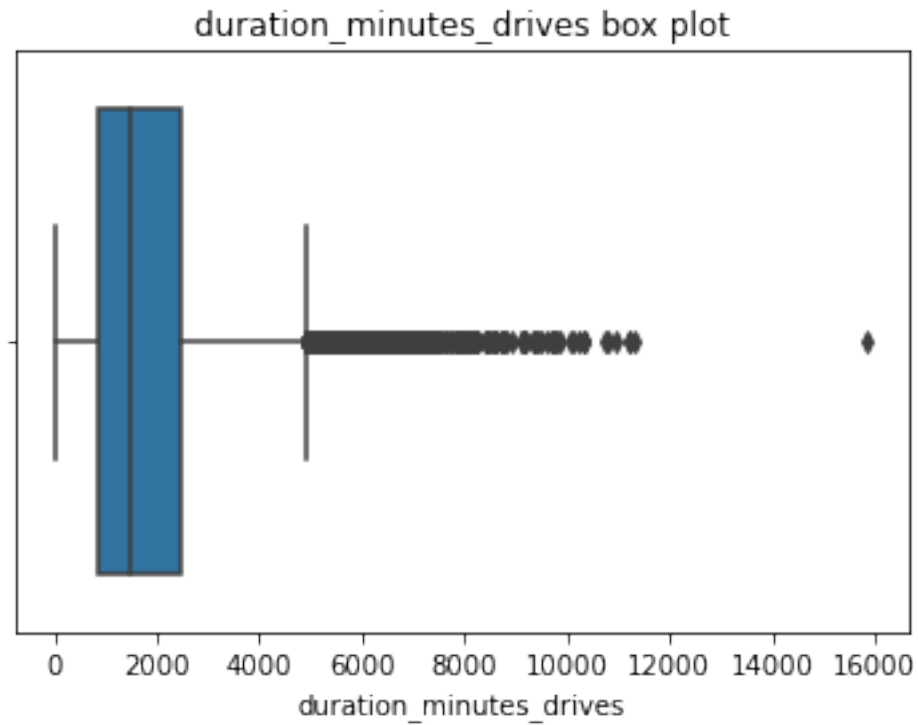


The number of drives driven in the last month per user is a right-skewed distribution with half the users driving under 3,495 kilometers. As you discovered in the analysis from the previous course, the users in this dataset drive *a lot*. The longest distance driven in the month was over half the circumference of the earth.

duration_minutes_drives *Total duration driven in minutes during the month*


```
[79]: # Box plot
sns.boxplot(x=df['duration_minutes_drives'])
plt.title('duration_minutes_drives box plot')
```

```
[79]: Text(0.5, 1.0, 'duration_minutes_drives box plot')
```



```
[81]: # Histogram
sns.histplot(x=df['duration_minutes_drives'])
median = df['duration_minutes_drives'].median()
plt.axvline(median, color='red', linestyle='-')
plt.text(2000,700,'median=1478.2', color='red')
plt.title('duration_minutes_drives Histogram')
```

```
[81]: Text(0.5, 1.0, 'duration_minutes_drives Histogram')
```

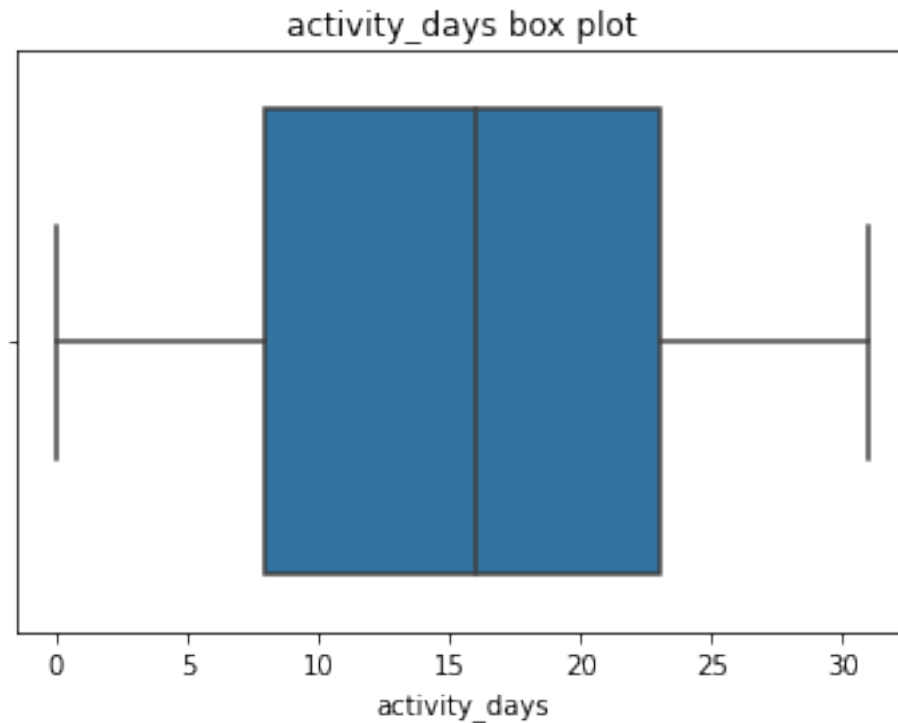



The `duration_minutes_drives` variable has a heavily skewed right tail. Half of the users drove less than ~1,478 minutes (~25 hours), but some users clocked over 250 hours over the month.

activity_days *Number of days the user opens the app during the month*

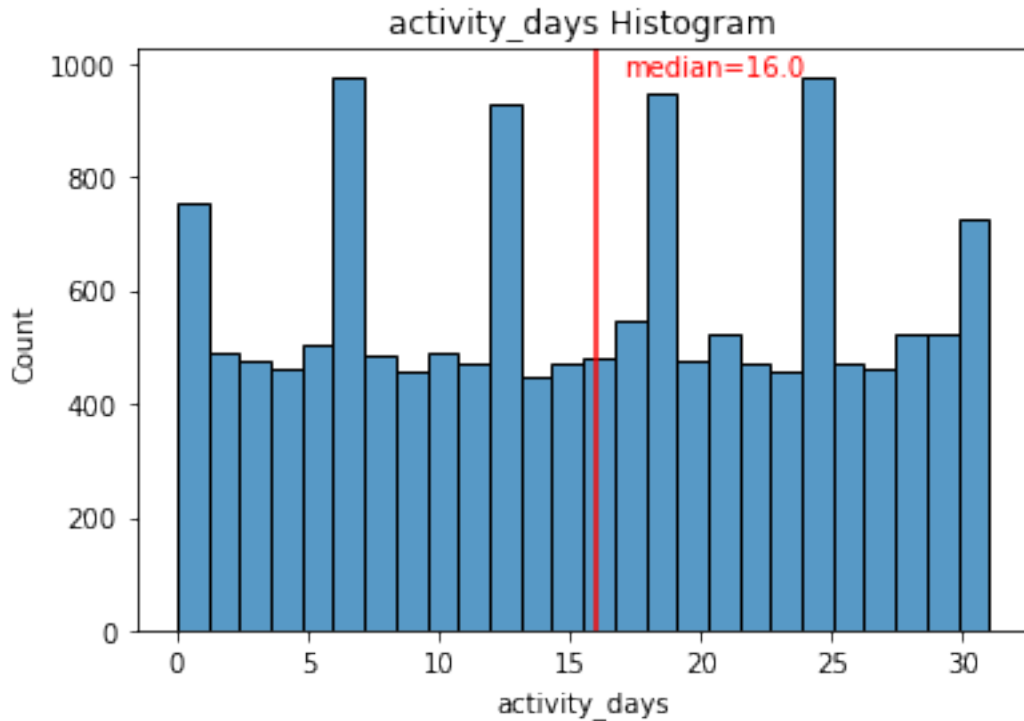
```
[82]: # Box plot
sns.boxplot(x=df['activity_days'])
plt.title('activity_days box plot')
```

```
[82]: Text(0.5, 1.0, 'activity_days box plot')
```



```
[92]: # Histogram
sns.histplot(x=df['activity_days'])
median = df['activity_days'].median()
plt.axvline(median, color='red', linestyle='-')
plt.text(17,980,'median=16.0', color='red')
plt.title('activity_days Histogram')
```

```
[92]: Text(0.5, 1.0, 'activity_days Histogram')
```




Within the last month, users opened the app a median of 16 times. The box plot reveals a centered distribution. The histogram shows a nearly uniform distribution of ~500 people opening the app on each count of days. However, there are ~250 people who didn't open the app at all and ~250 people who opened the app every day of the month.

This distribution is noteworthy because it does not mirror the `sessions` distribution, which you might think would be closely correlated with `activity_days`.

driving_days *Number of days the user drives (at least 1 km) during the month*


```
[84]: # Box plot
sns.boxplot(x=df['driving_days'])
plt.title('driving_days box plot')
```

```
[84]: Text(0.5, 1.0, 'driving_days box plot')
```



```
[88]: # Histogram
sns.histplot(x=df['driving_days'])
median = df['driving_days'].median()
plt.axvline(median, color='red', linestyle='--')
plt.text(13,1400,'median=12.0', color='red')
plt.title('driving_days Histogram')
```

```
[88]: Text(0.5, 1.0, 'driving_days Histogram')
```




The number of days users drove each month is almost uniform, and it largely correlates with the number of days they opened the app that month, except the `driving_days` distribution tails off on the right.

However, there were almost twice as many users (~1,000 vs. ~550) who did not drive at all during the month. This might seem counterintuitive when considered together with the information from `activity_days`. That variable had ~500 users opening the app on each of most of the day counts, but there were only ~250 users who did not open the app at all during the month and ~250 users who opened the app every day. Flag this for further investigation later.

device *The type of device a user starts a session with*

This is a categorical variable, so you do not plot a box plot for it. A good plot for a binary categorical variable is a pie chart.

```
[102]: # Pie chart
data=df['device'].value_counts()
plt.pie(data,
        labels=[f'{data.index[0]}: {data.values[0]}',
                f'{data.index[1]}: {data.values[1]}'],
        autopct='%1.1f%%')
plt.title('Users by device');
```




There are nearly twice as many iPhone users as Android users represented in this data.

label *Binary target variable (“retained” vs “churned”) for if a user has churned anytime during the course of the month*

This is also a categorical variable, and as such would not be plotted as a box plot. Plot a pie chart instead.

```
[103]: # Pie chart
data=df['label'].value_counts()
plt.pie(data,
        labels=[f'{data.index[0]}: {data.values[0]}',
                f'{data.index[1]}: {data.values[1]}'],
        autopct='%1.1f%%')
plt.title('Users by label');
```




Less than 18% of the users churned.

driving_days vs. activity_days Because both `driving_days` and `activity_days` represent counts of days over a month and they're also closely related, you can plot them together on a single histogram. This will help to better understand how they relate to each other without having to scroll back and forth comparing histograms in two different places.

Plot a histogram that, for each day, has a bar representing the counts of `driving_days` and `activity_days`.

```
[111]: # Histogram
label=['driving days', 'activity days']
plt.hist([df['driving_days'], df['activity_days']],
         label=label)
plt.legend()
plt.title('driving_days vs activity_days')
plt.xlabel('days')
plt.ylabel('count')
```

```
[111]: Text(0, 0.5, 'count')
```



As observed previously, this might seem counterintuitive. After all, why are there *fewer* people who didn't use the app at all during the month and *more* people who didn't drive at all during the month?

On the other hand, it could just be illustrative of the fact that, while these variables are related to each other, they're not the same. People probably just open the app more than they use the app to drive—perhaps to check drive times or route information, to update settings, or even just by mistake.

Nonetheless, it might be worthwhile to contact the data team at Waze to get more information about this, especially because it seems that the number of days in the month is not the same between variables.

Confirm the maximum number of days for each variable—`driving_days` and `activity_days`.

```
[113]: print(df['driving_days'].max())
print(df['activity_days'].max())
```

30


31

It's true. Although it's possible that not a single user drove all 31 days of the month, it's highly unlikely, considering there are 15,000 people represented in the dataset.

One other way to check the validity of these variables is to plot a simple scatter plot with the x-axis representing one variable and the y-axis representing the other.


```
[115]: # Scatter plot
sns.scatterplot(data=df, x='driving_days', y='activity_days')
plt.title('driving days vs activity days')
```

```
[115]: Text(0.5, 1.0, 'driving days vs activity days')
```




Notice that there is a theoretical limit. If you use the app to drive, then by definition it must count as a day-use as well. In other words, you cannot have more drive-days than activity-days. None of the samples in this data violate this rule, which is good.

Retention by device Plot a histogram that has four bars—one for each device-label combination—to show how many iPhone users were retained/churned and how many Android users were retained/churned.

```
[8]: # Histogram
sns.histplot(data=df,
             x='device',
             hue='label',
             multiple='dodge')
plt.title('Android vs iPhone retention histogram')
```

```
[8]: Text(0.5, 1.0, 'Android vs iPhone retention histogram')
```



The proportion of churned users to retained users is consistent between device types.

Retention by kilometers driven per driving day In the previous course, you discovered that the median distance driven last month for users who churned was 8.33 km, versus 3.36 km for people who did not churn. Examine this further.

1. Create a new column in `df` called `km_per_driving_day`, which represents the mean distance driven per driving day for each user.
2. Call the `describe()` method on the new column.

```
[9]: # 1. Create `km_per_driving_day` column
df['km_per_driving_day'] = df['driven_km_drives'] / df['driving_days']
# 2. Call `describe()` on the new column
df['km_per_driving_day'].describe()
```

```
[9]: count    1.499900e+04
mean          inf
std           NaN
min    3.022063e+00
25%    1.672804e+02
50%    3.231459e+02
75%    7.579257e+02
max          inf
```

Name: km_per_driving_day, dtype: float64

What do you notice? The mean value is infinity, the standard deviation is NaN, and the max value is infinity. Why do you think this is?

This is the result of there being values of zero in the `driving_days` column. Pandas imputes a value of infinity in the corresponding rows of the new column because division by zero is undefined.

1. Convert these values from infinity to zero. You can use `np.inf` to refer to a value of infinity.
2. Call `describe()` on the `km_per_driving_day` column to verify that it worked.

```
[10]: # 1. Convert infinite values to zero
df.loc[df['km_per_driving_day']==np.inf, 'km_per_driving_day'] = 0
# 2. Confirm that it worked
df['km_per_driving_day'].describe()
```


```
[10]: count      14999.000000
mean         578.963113
std          1030.094384
min           0.000000
25%          136.238895
50%          272.889272
75%          558.686918
max         15420.234110
Name: km_per_driving_day, dtype: float64
```

The maximum value is 15,420 kilometers *per drive day*. This is physically impossible. Driving 100 km/hour for 12 hours is 1,200 km. It's unlikely many people averaged more than this each day they drove, so, for now, disregard rows where the distance in this column is greater than 1,200 km.

Plot a histogram of the new `km_per_driving_day` column, disregarding those users with values greater than 1,200 km. Each bar should be the same length and have two colors, one color representing the percent of the users in that bar that churned and the other representing the percent that were retained. This can be done by setting the `multiple` parameter of seaborn's `histplot()` function to `fill`.

```
[12]: # Histogram
plt.figure(figsize=(10,5))
sns.histplot(data=df,
             x='km_per_driving_day',
             bins=range(0,1300,30),
             hue='label',
             multiple='fill')
plt.ylabel('%')
plt.title('Churn rate by average km per driving day')
```

```
[12]: Text(0.5, 1.0, 'Churn rate by average km per driving day')
```




The churn rate tends to increase as the mean daily distance driven increases, confirming what was found in the previous course. It would be worth investigating further the reasons for long-distance users to discontinue using the app.

Churn rate per number of driving days Create another histogram just like the previous one, only this time it should represent the churn rate for each number of driving days.

```
[16]: # Histogram
plt.figure(figsize=(10,5))
sns.histplot(data=df,
             x='driving_days',
             hue='label',
             multiple='fill')
plt.ylabel('%')
plt.title('Churn rate per driving day')
```

```
[16]: Text(0.5, 1.0, 'Churn rate per driving day')
```



The churn rate is highest for people who didn't use Waze much during the last month. The more times they used the app, the less likely they were to churn. While 40% of the users who didn't use the app at all last month churned, nobody who used the app 30 days churned.

This isn't surprising. If people who used the app a lot churned, it would likely indicate dissatisfaction. When people who don't use the app churn, it might be the result of dissatisfaction in the past, or it might be indicative of a lesser need for a navigational app. Maybe they moved to a city with good public transportation and don't need to drive anymore.

Proportion of sessions that occurred in the last month Create a new column `percent_sessions_in_last_month` that represents the percentage of each user's total sessions that were logged in their last month of use.

```
[17]: df['percent_sessions_in_last_month'] = df['sessions'] / df['total_sessions']
```

What is the median value of the new column?


```
[18]: df['percent_sessions_in_last_month'].median()
```

```
[18]: 0.42309702992763176
```

Now, create a histogram depicting the distribution of values in this new column.

```
[26]: # Histogram
sns.histplot(x=df['percent_sessions_in_last_month'],
             hue = df['label'])
median = df['percent_sessions_in_last_month'].median()
plt.axvline(median, color='red', linestyle='--')
```

[26]: <matplotlib.lines.Line2D at 0x7fd77f075050>



Check the median value of the `n_days_after_onboarding` variable.

```
[27]: df['n_days_after_onboarding'].median()
```

[27]: 1741.0


Half of the people in the dataset had 40% or more of their sessions in just the last month, yet the overall median time since onboarding is almost five years.

Make a histogram of `n_days_after_onboarding` for just the people who had 40% or more of their total sessions in the last month.

```
[28]: # Histogram
data = df.loc[df['percent_sessions_in_last_month'] >= 0.4]
sns.histplot(x=data['n_days_after_onboarding'])
plt.title('Days after onboarding for users with more than 40% sessions last_
↪month')
```

```
[28]: Text(0.5, 1.0, 'Days after onboarding for users with more than 40% sessions last
month')
```

Days after onboarding for users with more than 40% sessions last month



The number of days since onboarding for users with 40% or more of their total sessions occurring in just the last month is a uniform distribution. This is very strange. It's worth asking Waze why so many long-time users suddenly used the app so much in the last month.

5.0.1 Task 3b. Handling outliers

The box plots from the previous section indicated that many of these variables have outliers. These outliers do not seem to be data entry errors; they are present because of the right-skewed distributions.

Depending on what you'll be doing with this data, it may be useful to impute outlying data with more reasonable values. One way of performing this imputation is to set a threshold based on a percentile of the distribution.

To practice this technique, write a function that calculates the 95th percentile of a given column, then imputes values > the 95th percentile with the value at the 95th percentile. such as the 95th percentile of the distribution.

```
[35]: def outlier_imputer(column_name, percentile):
        threshold = df[column_name].quantile(percentile)
        df.loc[df[column_name] > threshold, column_name] = threshold

        print('{>25} | percentile: {} | threshold: {}'.format(column_name,
        ↪percentile, threshold))
```

Next, apply that function to the following columns: * sessions * drives * total_sessions * driven_km_drives * duration_minutes_drives

```
[36]: for column in ['sessions', 'drives', 'total_sessions',
                    'driven_km_drives', 'duration_minutes_drives']:
        outlier_imputer(column, 0.95)
```

```

sessions | percentile: 0.95 | threshold: 243.0
drives   | percentile: 0.95 | threshold: 201.0
total_sessions | percentile: 0.95 | threshold: 454.3415426984
driven_km_drives | percentile: 0.95 | threshold: 8889.776003166
duration_minutes_drives | percentile: 0.95 | threshold: 4668.8139585300005
```

Call describe() to see if your change worked.

```
[37]: df.describe()
```

```
[37]:
```

	ID	sessions	drives	total_sessions \
count	14999.000000	14999.000000	14999.000000	14999.000000
mean	7499.000000	76.568705	64.058204	184.030237
std	4329.982679	67.297958	55.306924	118.597994
min	0.000000	0.000000	0.000000	0.220211
25%	3749.500000	23.000000	20.000000	90.661156
50%	7499.000000	56.000000	48.000000	159.568115
75%	11248.500000	112.000000	93.000000	254.192341
max	14998.000000	243.000000	201.000000	454.341543

	n_days_after_onboarding	total_navigations_fav1 \
count	14999.000000	14999.000000
mean	1749.837789	121.605974
std	1008.513876	148.121544
min	4.000000	0.000000
25%	878.000000	9.000000
50%	1741.000000	71.000000
75%	2623.500000	178.000000
max	3500.000000	1236.000000

	total_navigations_fav2	driven_km_drives	duration_minutes_drives \
count	14999.000000	14999.000000	14999.000000
mean	29.672512	3939.631852	1789.643156
std	45.394651	2216.039474	1222.695112
min	0.000000	60.441250	18.282082
25%	0.000000	2212.600607	835.996260
50%	9.000000	3493.858085	1478.249859
75%	43.000000	5289.861262	2464.362632
max	415.000000	8889.776003	4668.813959

	activity_days	driving_days	km_per_driving_day \
--	---------------	--------------	----------------------

count	14999.000000	14999.000000	14999.000000
mean	15.537102	12.179879	578.963113
std	9.004655	7.824036	1030.094384
min	0.000000	0.000000	0.000000
25%	8.000000	5.000000	136.238895
50%	16.000000	12.000000	272.889272
75%	23.000000	19.000000	558.686918
max	31.000000	30.000000	15420.234110

	percent_sessions_in_last_month
count	14999.000000
mean	0.449255
std	0.286919
min	0.000000
25%	0.196221
50%	0.423097
75%	0.687216
max	1.530637

Conclusion Analysis revealed that the overall churn rate is ~17%, and that this rate is consistent between iPhone users and Android users.

Perhaps you feel that the more deeply you explore the data, the more questions arise. This is not uncommon! In this case, it's worth asking the Waze data team why so many users used the app so much in just the last month.

Also, EDA has revealed that users who drive very long distances on their driving days are *more* likely to churn, but users who drive more often are *less* likely to churn. The reason for this discrepancy is an opportunity for further investigation, and it would be something else to ask the Waze data team about.

5.1 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

5.1.1 Task 4a. Results and evaluation

Having built visualizations in Python, what have you learned about the dataset? What other questions have your visualizations uncovered that you should pursue?

Pro tip: Put yourself in your client's perspective. What would they want to know?

Use the following code fields to pursue any additional EDA based on the visualizations you've already plotted. Also use the space to make sure your visualizations are clean, easily understandable, and accessible.

Ask yourself: Did you consider color, contrast, emphasis, and labeling?

==> ENTER YOUR RESPONSE HERE

I have learned that: Number of Iphone users for Waze app nearly double number of Android users Around 18% of users churned For most of days in a month, there are more people who drive than using Waze app, but there are also few days when there are more people using Waze app. Driving days and activity days have strong positive correlation. Proportion of churned users between Iphone and Android devices is consistent The more average km per day users drive, the higher the churning rate

My other questions are Why does 'label' columns have quites lots of missing rows? Why retained users have fewer driving days?

My client would likely want to know ... Which variable/factor directly affect user churning rate? How can they reduce user churning rate?

5.1.2 Task 4b. Conclusion

Now that you've explored and visualized your data, the next step is to share your findings with Harriet Hadzic, Waze's Director of Data Analysis. Consider the following questions as you prepare to write your executive summary. Think about key points you may want to share with the team, and what information is most relevant to the user churn project.

Questions:

1. What types of distributions did you notice in the variables? What did this tell you about the data?
 2. Was there anything that led you to believe the data was erroneous or problematic in any way?
 3. Did your investigation give rise to further questions that you would like to explore or ask the Waze team about?
 4. What percentage of users churned and what percentage were retained?
 5. What factors correlated with user churn? How?
 6. Did newer uses have greater representation in this dataset than users with longer tenure? How do you know?
-
1. Most variables have right skewed and uniform distributions. For right skewed distribution, it means that there are more users with the lower values. For the uniform distribution, it means value varies and dont have a pattern.
 2. It seems that nothing is completely wrong about the data or any variables. There is just one concern about difference of maximum days of 'driving days' and 'activity days' although this may not cause significant change in analysis.
 3. Why did so many people(40% users) suddenly use Waze app last month?
 4. Around 18% users churned while 82% retained
 5. km_per_driving_day has positive correlation with churn because the more km people drive a day, the higher the churning rate. On the other hand, number of driving days has negative correlation with churning rate. The more days people drive, the lower the churning rate.

6. Based on `n_days_after_onboarding` histogram, we can see that the answer here is no because the distribution is uniform.

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.